

Free Speech as Causal Transparency

Abstract

Free speech is invoked constantly in political and legal discourse, but there is no widely agreed account of what makes speech free. We propose a causal-transparency account of free speech, organised by a taxonomy of speech constraints that we adapt from Lowe but motivate independently. An utterance is freely spoken, on our account, when the causal pathway from the speaker’s communicative intentions to the utterance is recoverable by a competent interlocutor given the words and the shared context. The degree of freedom is inversely proportional to the cost the speaker faces in making this pathway transparent. The account is stated over the full space of possible utterances available to a speaker on a topic in a context, including silence, and parametrised by topic-context pairs rather than treated as a global binary property. Key consequences include: lying is excluded from free speech on structural grounds independent of moral evaluation, tactical speech is free, coercion and censorship reduce degree of freedom on a common scale, and the propagation of suppression through inferentially connected topics is analytically tractable. The framework also gives formal content to Bejan’s distinction between isegoria and parrhesia, separates degree of freedom from a speaker’s willingness to reveal, and interprets free speech law. We extend Lowe’s taxonomy with a Type 0 category for obstructions at the level of thought-formation, and develop the account in conversation with work across philosophy of language, speech-act theory, and the history of political thought.

Word count: 9,204 words, including notes and references.

1 Introduction

Debates about free speech are rarely about whether speech should be free. They are about which constraints on speech count as restrictions of freedom and which do not. A government that punishes dissent restricts free speech, while a government that prohibits defamation, on most views, does not. But the line is drawn by intuition and precedent rather than by any positive account of free speech. Rebecca Lowe’s proposed taxonomy (Lowe 2026) distinguishes three types of obstruction: Type 1, which concerns the speaker’s capacity to produce an utterance and determine its content, Type 2, which concerns the communicative channel connecting speaker to interlocutor, and Type 3, which presupposes both but involves the speaker refraining due to perceived risk of

adverse consequences. For example, censorship, coercion, and chilling effects are all Type 3. The taxonomy is useful as it allows us to specify which kind of obstruction is operative before arguing about entitlement. However, Lowe’s characterisation of Type 3 is negative in form: the presence of a threat or anticipated cost marks an utterance as unfree, and no condition is given under which speech counts as free. The positive account of what free speech consists in is left open.

This paper offers that positive account, which we call the causal-transparency condition. The account is parametrised from the outset: freedom of speech is assessed relative to a topic and a context, as a field over topic-context pairs rather than as a global property of a speaker. Our main contributions are as follows.

1. **The causal-transparency condition** (section 4). We give positive content to Type 3 free speech: an utterance is freely spoken when the causal pathway from the speaker’s communicative intentions to the utterance is recoverable by a competent interlocutor. The condition is stated over the full space of possible utterances $U(S, T, C)$ available to a speaker S on a topic T in a context C , including silence. Whether opacity is self-engineered or externally imposed is irrelevant, as only recoverability matters. This places lying and coercion on a common analytical axis. We argue that lying is a symptom of constrained expression: a speaker who lies to achieve goal G cannot achieve it through transparent expression.
2. **Type 0** (sections 2.1 and 3). We extend Lowe’s taxonomy with a Type 0 category: obstructions at the level of thought-formation, as an upstream background condition that frames the scope of the analysis.
3. **Formal content for Bejan’s isegoria/parrhesia** (section 5.1). Parrhesia maps onto the within-cell causal-transparency condition. Isegoria is identified with the conjunction of Types 0, 1, and 2: having a thought one can form, express, and transmit to a reachable interlocutor.
4. **Suppression topology** (section 5.2). Because freedom is parametrised over topic-context (T, C) pairs, suppression in one cell propagates along inferential connections to adjacent cells, making the topology of constrained speech analytically tractable.
5. **Scalar degree of freedom and willingness-to-reveal** (sections 4.2 and 5.3). Freedom admits of degree, measured inversely by the cost the speaker faces in making the pathway transparent. This handles chilling effects and the full range of intermediate cases that binary accounts cannot distinguish. We further identify willingness-to-reveal as a distinct parrhesiastic property, best studied via a network of contexts adjacent to the original (T, C) cell.
6. **A reading of free speech law** (section 5.4). Legal free speech protection is interpreted using our framework as applicable only in specific contexts.

2 Related Work

2.1 Lowe’s Taxonomy of Speech Constraints

Rebecca Lowe’s “Speaking Freely” project (Lowe 2026) offers the most useful structural starting point. We draw on Lowe’s taxonomy as presented in a podcast discussion and summarised on her Substack. This work is in progress and no peer-reviewed version is yet available. Lowe proposes that obstructions to free speech fall into three hierarchically related types, and that public discourse on free speech would be substantially clarified by specifying which type is operative in any contested case before arguing about entitlement.

Type 1 obstructions concern the speaker’s physical or cognitive capacity: laryngitis, a hand over the mouth, or any external agent directly controlling the content of the speaker’s utterance. Type 2 obstructions presuppose Type 1 capacity but concern the communicative channel: the speaker can speak and control what they say, but the channel connecting them to their intended interlocutor is blocked. Type 3 obstructions presuppose both: the speaker has full physical capacity and an open channel, but refrains due to perceived risk of social, physical, or psychological consequences. Censorship and coercion are paradigmatic Type 3 cases, and so, on Lowe’s account, is self-censorship that follows from anticipating social cost.

The taxonomy is diagnostically useful but deliberately incomplete at Type 3. Lowe characterises it negatively: the presence of a threat or anticipated cost is what marks speech as unfree, and no condition is given for when speech is free. Our account supplies the converse, positive direction: where the causal pathway from intention to utterance is reconstructible, speech is free. This is the gap we address here with the starting assumption that Types 1 and 2 constraints are satisfied. Finally, we offer a Type 0 category of obstructions, which has to do with the ability to formulate certain thoughts in the first place.

We subscribe to Lowe’s three-way division essentially verbatim. We do not rely on Lowe’s authority for the taxonomy, and motivate it independently (including our own Type 0, and our Type 3 definition refinement). A communicative intention must first be formed, and its obstruction is what we call Type 0. It is then expressed in an utterance that fixes its content (Type 1), then carried through a channel to an interlocutor (Type 2), and finally recovered by that interlocutor from the utterance and the shared context (Type 3). Each stage is a distinct point of possible failure, and the failures are of correspondingly distinct kinds. We take from Lowe the diagnostic insight that separating these stages clarifies disputes about free speech.

That speech constraints can be typologised rigorously is not, in any case, peculiar to Lowe. Cook and Heilmann (2013) give a peer-reviewed typology of a neighbouring phenomenon, distinguishing public self-censorship, in which an agent suppresses expression in response to an external censor, from private self-censorship, in which censor and censee coincide and no external censor is

present. Their axis differs from Lowe’s, being the identity of the censor rather than the stage of the communicative process obstructed, but it sharpens what Lowe files under Type 3, where self-censorship in the face of anticipated cost is treated as one thing. They draw an implication that principles of free speech do not apply directly to private self-censorship, since absent an external censor the suppression is non-coercive. This is implied by our account: since degree of freedom is fixed by the cost of transparency imposed on the speaker, self-restraint in the absence of any external cost leaves freedom untouched, however much it reduces what is said. Private self-censorship (at least above the involuntary Type 0 restriction) is thus, on our framework as on theirs, not a matter for free speech. We return to the analytical/remedial separation this reflects in section 6.3.

2.2 Shiffrin’s Thinker-Based Account

The closest existing account to the one we develop is Seana Shiffrin’s thinker-based theory of free speech, set out in *Speech Matters* (Shiffrin 2014). Shiffrin grounds the scope and structure of communicative liberty in what is necessary for the meaningful exercise of our capacities as thinking beings. Among the conditions required for beings to function as thinkers and moral agents is the existence of reliable channels through which speakers can convey their actual mental contents to one another. Sincere transmission of genuine beliefs is more than just a use to which communication is sometimes put. It is what makes communication a resource capable of sustaining the activities (promising, consenting, confessing, disagreeing) that require others to know what one actually thinks.

From this, Shiffrin derives what she takes to be an internal limit on free speech: lying falls outside its protective ambit. A lie does not just harm the listener, it transforms a mechanism for conveying the speaker’s genuine mental contents into one that can convey either true or false representations of those contents, corrupting the infrastructure on which free speech’s value depends. Protection of sincere expression and prohibition of lying share the same foundation. The thinker-based approach extends symmetrically to compelled speech (Shiffrin 2011): what is most objectionable about forcing a speaker to say things they do not believe is that it interferes with the autonomous thought processes of the compelled speaker, rather than the listener’s ability to recover them.

We share Shiffrin’s central commitment: free speech is grounded in the fidelity of the pathway from the speaker’s mental states to their utterance, and lying falls outside free speech for this reason. Our account diverges from hers at three points.

First, Shiffrin treats lying and coercion as structurally distinct violations: lying manipulates the epistemic environment of the listener, while coercion violates the expressive autonomy of the speaker. We argue that the source of opacity (whether self-engineered or externally imposed) is irrelevant to the freedom condition. Only recoverability matters. This unification reflects a deeper claim

we develop in section 4: lying is itself a form of constrained expression. Shiffrin’s account lacks this positive account of why lying is constrained rather than merely corrupting: for her, lying is excluded because it corrupts a communicative resource that moral agency requires, not because it manifests a constraint on the liar’s own expression.

Second, Shiffrin’s account is binary. We argue in section 5.3 that freedom of speech admits of degree, measured by the cost to the speaker of making the causal pathway transparent. This scalar extension is required to handle chilling effects and the full range of intermediate cases.

Third, Shiffrin does not parametrise free speech by topic and context. We argue that freedom is a field over topic-context pairs: a speaker may be fully free with respect to some topics in some contexts while substantially constrained with respect to others. This changes the structure of the analysis, since it allows censorship to be characterised as a local disruption of specific cells rather than a global condition, and makes the topology of topic-space analytically tractable.

2.3 Grice, Speech Act Theory, and the Recoverability Condition

Paul Grice’s account of speaker’s meaning provides the technical foundation for the recoverability condition (Grice 1957, 1975). For Grice, natural meaning is the non-intentional relation between symptom and cause. Non-natural meaning arises when a speaker intends an effect on an audience, intends that this intention be recognised, and intends that the recognition be part of how the effect comes about. It is this reflexive communicative intention that we are asking to be recoverable: free speech, on our account, is speech in which a competent interlocutor can reconstruct this chain from words and shared context.

Grice’s quality maxim, which states “do not say what you believe to be false”, is where his cooperative principle most directly intersects with our framework. The relevant observation is not the normative one that “lying violates cooperative norms”, but the structural one: a lie is an utterance whose causal pathway from the speaker’s mental state cannot be recovered using the inferential tools that conversational cooperation makes available. The interlocutor who applies Grice’s reasoning to a sincere utterance can recover the speaker’s meaning, but the same tools applied to a lie yield a false result.

Post-Gricean work has refined the inferential picture without changing its overall shape. Sperber and Wilson’s relevance theory (Sperber and Wilson 1995) grounds inferential comprehension in a cognitive principle of relevance: the interpretation a hearer settles on is the one that delivers the most cognitive effect for the least processing cost. This provides an empirically grounded account of how recovery actually proceeds in practice: a competent interlocutor’s reconstruction of the causal pathway is what a relevance-driven inferential process produces under the shared context. Relevance theory supplies a cognitively realistic mechanism for the recovery condition.

Rae Langton and Jennifer Hornsby’s work on illocutionary silencing (Langton 1993; Hornsby and Langton 1998) develops a related analysis from the listener side. Drawing on Austin’s taxonomy of speech acts, they argue that a speaker’s illocutionary act depends on uptake: the hearer’s recognition of the speaker’s communicative intention. When the communicative context is structured so that uptake is systematically withheld from certain speakers, their illocutionary acts misfire. The locutionary act is performed, but the illocutionary act fails. Langton and Hornsby’s uptake condition is the recoverability condition viewed from the listener’s side. Their work shows that the success of a speech act cannot be read off the speaker’s act alone: it depends on whether the causal pathway from the speaker’s intention to the interlocutor’s recognition is traversable within the context. Our definition makes this speaker-interlocutor-context structure explicit and turns it into a criterion for freedom.

Mary Kate McGowan (2019) extends the Langton-Hornsby framework by analysing how individual utterances shift the conversational norms that govern subsequent moves, “permitting” certain follow-ups by setting factual presumption, conversational tone, or normative defaults. This bears directly on the inferential-propagation argument we develop in section 5.2: constraints on speech in one (T, C) cell spread to adjacent cells not only through inferential implication but through the conversational norms that prior utterances institute. McGowan’s machinery supplies a finer-grained account of how suppression in one cell can deform the conditions for free speech in connected cells.

The framework relates more broadly to the Austin-Searle tradition of speech act theory (Austin 1962; Searle 1969). Recoverability applies uniformly across Searle’s speech act types (assertives, directives, commissives, expressives, and declarations) since what is being recovered is the pathway from the speaker’s communicative intention to the utterance, whatever specific act that intention encodes. Sincerity conditions on speech acts, central to Searle’s account, become structural rather than normative on our reading: the violation of sincerity in lying is what renders the causal pathway opaque. Coerced performatives (forced confessions, mandatory recitations, ceremonial utterances under threat) fall outside the free speech region for the same structural reason as the coerced assertives discussed in section 4.3: the causal origin of the utterance has been replaced, regardless of which type of act is nominally performed.

To our knowledge, no major account in the free speech literature has systematically used Grice’s conditions on the recovery of speaker’s meaning as the basis for defining what makes speech free.

2.4 Bejan on Isegoria and Parrhesia

Teresa Bejan’s work (Bejan 2017, 2019, 2020) recovers a distinction from ancient Greek political thought that contemporary free speech discourse has largely collapsed. Isegoria designated the equal right of citizens in good standing to address the popular assembly and receive a hearing. It was an institutional

concept tied to democratic membership and equal formal access. Parrhesia designated the licence to speak frankly, to say what one actually thinks even at personal cost. It required both the speaker's courage and the audience's tolerance in the classical sense of enduring what one dislikes without suppressing it.

For our framework, Bejan's distinction provides vocabulary for two properties that our analysis can articulate in terms of a single definition of free speech. Parrhesia maps onto the causal-transparency condition and the willingness-to-reveal property we discuss in sections 5.1 and 5.3: parrhesiastic speech is speech in which the pathway from the speaker's genuine beliefs to their utterance is intact and the speaker would make it transparent even at cost. Isegoria maps onto a structural precondition for the topic-context parametrisation: which cells are available to a speaker is a question about equal access to the kinds of communicative situations in which speech can occur, including which interlocutors are reachable and what authority the context confers on the speaker's utterances. A speaker may be fully parrhesiastic within the cells available to them while lacking isegoria-style access to cells where their speech would carry weight.

2.5 Mill, Schauer, and the Broader Free Speech Tradition

The dominant historical argument for free speech is Mill's (1859): protect free expression because doing so is the most reliable route to truth. A silenced opinion may be true (truth is lost), partially true (partial truth is lost), or wholly false (in which case unsilenced opposition keeps the received opinion from decaying into dead dogma, on Mill's view). The justification is consequentialist about epistemic outcomes: free discussion is valuable because of what it produces.

Our account differs from Mill's. The causal-transparency condition does not require that the speaker's beliefs be true, well-formed, or rationally grounded. Truth-tracking enters our framework downstream rather than upstream: when speakers can transparently convey their mental states, the social processes that select for better beliefs (deliberation, criticism, error-correction) have the inputs they need. Mill's argument can be read as one substantive specification of what those processes produce. We make no such commitment, demonstrating how a theory of free speech can be given without taking on a controversial empirical claim about what that discussion delivers.

Schauer (1982) frames the foundational question slightly differently, asking why free speech requires special protection beyond general liberty. He surveys arguments from truth (Mill), from democracy, and from autonomy, and concludes that no single justification grounds the strong protections free speech is typically accorded. Our account is closest in spirit to the autonomy strand but reformulates it structurally rather than substantively. The recoverability condition picks out exactly what is at stake in communicative liberty (the integrity of the pathway from mental state to expression), which general liberty principles do not capture. A regime could permit speech, in the sense of not intervening at the act of

utterance, while systematically undermining recoverability through control of context, channel, or vocabulary. Our account makes the distinction explicit and gives it analytical content.

On Scanlon’s (1972) early view, a legitimate government may not restrict expression on the ground that citizens might be persuaded by it to form false beliefs or harmful intentions, because doing so is incompatible with their standing as autonomous agents who must weigh reasons for themselves. Free speech protects the audience’s sovereignty over its own mind. In our account, this can be expressed as protection against Type 0 constraints. However, our Type 3 allows this situation to be viewed from the perspective of the speaker and the recoverability of their communicative pathway, an interest prior to and independent of what the hearer then does with what they recover. The two are complementary. Recoverability is the condition under which the autonomy Scanlon prizes has an object.

A separate literature asks whether freedom is the kind of thing that admits of degree at all, and how overall freedom might be measured (Kramer 2002, 2003; Carter 1999). We take up its bearing on our scalar claim in section 6.2.

3 Simplifying Assumptions

We restrict the scope of the analysis to intentional speech acts. Slips, compelled utterances arising from neurological causes, and other cases where speaking is not under the agent’s control fall outside the domain we are analysing.

Unless otherwise stated, we treat Lowe’s Types 1 and 2 not as obstacle categories but as background conditions shared between speaker and interlocutor. Specifically, we assume: (a) the speaker has the physical and cognitive capacity to speak and to determine the content of their utterance, and is a reasonable agent (in the sense that whatever they utter issues from some communicative reason rather than being produced at random), (b) a functional communicative channel and a shared language exist between speaker and interlocutor, and (c) the information environment shaping the speaker’s belief-formation is mutually known to both parties, though not required to be accurate. These correspond to Lowe’s Types 1, 2, and the upstream condition we call Type 0 respectively.

Condition (c) requires epistemic symmetry about the speaker’s situation, not accuracy. A speaker whose beliefs have been shaped by a distorted information environment satisfies (c) if the interlocutor’s beliefs admit a similar distortion. The distortion is part of the shared context, used in recovering the causal pathway.

The conditions are not fully independent, and government intervention can target each level distinctly. Persecuting the use of minority languages operates at Type 1: it severs the ability of one generation to speak to the next in their own language, removing expressive capacity even when speakers remain cognitively capable of forming the relevant thoughts. Such persecution can also cause upstream effects

at Type 0 over time: a language denied intergenerational transmission eventually ceases to be a living conceptual resource, removing with it the substrate through which certain thoughts are most naturally formed. Newspeak-style vocabulary restriction is a milder instance of the same dynamic: constraining expression at Type 1 with potential long-run Type 0 consequences. At Type 2, government control of communicative infrastructure severs or filters channels before any question of cost on expression arises. All three forms of intervention foreclose the preconditions our analysis of Type 3 freedom takes as given rather than operating within them, and a full treatment lies outside the scope of this paper.

In section 5.3 we treat the cost of making the causal pathway transparent as a single scalar quantity, an idealisation we call linearisation: it places plural (and possibly incommensurable) costs on one common linear scale. This is a significant simplification. The costs a speaker faces may be incommensurable: social harm to others, personal risk, and inconvenience are different kinds of cost with no obvious common scale. Costs are also speaker-relative: the same consequence carries different weight for different agents. We adopt this linearisation as a useful idealisation that captures the right qualitative structure without claiming metric precision. We consider this an acceptable simplification for the reason that in practice, the scale is outwardly simplified to binary: the speaker either shares or does not share a given thought in a given context.

We further assume: (d) the speaker can assess the cost structure of speaking in the given context, and (e) the speaker has an accurate model of what a competent interlocutor knows, can understand, and can conclude from the shared context. The “competent interlocutor” standard is fixed by this: an interlocutor with the cognitive and contextual capacities the speaker correctly attributes to them under (e). Competence is therefore not a free parameter floating above the analysis but is anchored in the speaker’s accurate model. Agents who lack (d) because they are unfamiliar with the relevant social or political norms, or have relevant cognitive limitations, present a case outside the framework as stated. The symmetric case, where either party’s interpretation of the shared context is mistaken, e.g. secret surveillance scenarios, is likewise outside scope.

4 The Definition

4.1 The Space of Possible Utterances

We analyse free speech not as a property of individual past utterances but as a property of the space of possible utterances available (in the sense of lack of Type 0-2 obstacles) to an agent. Given assumptions (a)-(e) of section 3, we fix a speaker S , a topic T , and a context C , and consider the full set $U(S, T, C)$ of utterances S could produce on T in C . This includes silence: the decision not to speak is itself an element of $U(S, T, C)$, and the question of whether S is free to remain silent falls within the same framework.

On the reasonable-agent assumption (a) of section 3, a speaker making an utterance has some reason for doing so. However, no agent is ever free over

the entire space $U(S, T, C)$ of the utterances they have reasons for considering speaking. The vast majority of one's thoughts are not something they wish to express. Lying to produce a desired outcome (section 4.3) is the salient case. Many other utterances are sincere and recoverable but carry costs. The question of free speech is therefore not whether S can say anything, but which subset of $U(S, T, C)$ is freely accessible and to what degree.

Freedom of speech, so understood, is not a binary global property of a speaker but a field over (T, C) pairs: a speaker may be fully free with respect to some topics in some contexts while substantially constrained with respect to others. The causal-transparency condition introduced next operates within a single cell, assessing which elements of $U(S, T, C)$ are freely spoken. Claims about a speaker's freedom more generally range over the cells themselves. T and C are not formally distinguished in the present analysis: the pair (T, C) is treated as a unit, and the topic parameter is a convenience for picking out cases where a specific subject matter is salient (political affairs, personal life, professional conduct) rather than a structurally necessary component of the framework.

4.2 The Causal-Transparency Condition

Let us assume the simplifying assumptions in section 3 hold. Then, we say that an utterance u in $U(S, T, C)$ is freely spoken if the causal pathway from S 's mental states to u is recoverable by a competent interlocutor given the words and the shared context. More precisely, u is freely spoken to the extent that a competent interlocutor, drawing on the utterance and the shared context, can reconstruct the chain from S 's mental states to u .

The mental states in question are not S 's general psychological states but specifically those which form S 's communicative intention on T : the beliefs, desires, and intentions that are, roughly, necessary and sufficient (in some simplified sense) for the speaker to formulate u . Background states are relevant only insofar as they bear on this intention.

This is a two-stage assessment. First, recoverability is a binary prerequisite: if the pathway is opaque, so that no competent interlocutor could reconstruct it from the available information, the utterance falls outside the free speech region of $U(S, T, C)$ entirely. We treat this as yet another simplifying idealisation: in practice, pathways may be partially recoverable, i.e. an interlocutor may be able to tell that S is being constrained without knowing exactly what S would otherwise say. Modelling partial recoverability would require a graded notion, so we set it aside as it does not change the structure of the analysis. Second, within the recoverable region, degree of freedom varies: an utterance is more freely spoken the lower the cost to S of making the pathway transparent, where cost includes the anticipated social, psychological, and physical consequences of doing so, subject to the linearisation assumption of section 3.

The source of opacity is irrelevant: whether the pathway is influenced by an external agent, by fear of consequences, or by S 's own design, an unrecoverable

pathway is not freely spoken.

4.3 Key Cases

Coercion. Coercion at gunpoint admits two readings within the framework, and it is instructive that both are available. On the first, the causal origin of the utterance has been replaced: what S says is caused by the coercer’s will rather than by S’s mental states, leaving no pathway from S’s mind for an interlocutor to recover. So construed, coercion approximates a Type 1 obstruction, since the capacity to determine content has been overridden, even though the channel is open and S can physically speak. On the second, the pathway is not eliminated but made prohibitively costly: S retains the capacity to say what they think and complies because the cost of transparency is extreme. So construed, coercion is the limiting case of the chilling effects analysed below, on the common scale of section 5.3. Some people do tell the truth at gunpoint, at great cost, which shows that the pathway and the capacity typically remain intact, and the second reading is therefore the more general one, with the first recovered as its limit as the cost of transparency grows without bound. On either reading the utterance is not freely spoken: conditions (a) and (b) are satisfied, but the pathway is either absent or recoverable only at a cost that drives the degree of freedom to near zero.

Censorship and chilling effects. A state that punishes dissent does not block the channel or remove physical capacity. Instead, it raises the cost of transparency in specific (T, C) cells while leaving others intact. The anticipated cost of sincerity in constrained cells deforms the causal pathway: what S says is determined partly by fear as part of S’s mental states. A competent interlocutor who knows the censorship context can often recover the gap between utterance and belief. This is why speech under obvious censorship is frequently read correctly by all parties, but direct speech remains unfree to the degree that the cost of transparency is high. The rise of the use of the term “unalive” is an example. What such a state restricts is thus a targeted subset of the topic-context space, not speech as such: censorship is a local disruption of the causal-transparency condition. How this local disruption spreads to adjacent cells is taken up in section 5.2.

Tactical speech. A speaker who declines a cookie before dinner, saying “I don’t want one” when they do, is exercising self-control. Whether this is freely spoken depends on context: if a competent interlocutor can recover from the situation that the speaker is likely exercising self-restraint rather than expressing a genuine preference, the pathway is recoverable and the utterance is freely spoken. The speaker is not engineering opacity because the gap between utterance and mental state is available to be inferred.

Silence and compelled speech. Silence is an element of $U(S, T, C)$, so the framework covers it directly. A speaker who chooses silence speaks freely when a competent interlocutor can recover why nothing was said, i.e. that

the silence issues from the speaker’s own communicative intention rather than from an imposed cost. Compelled speech is the mirror image, since forcing a speaker to utter what they would otherwise leave unsaid removes silence from the freely accessible set, or raises the cost of remaining silent, in the same way that censorship raises the cost of speaking. The Soviet voting example of section 5.3 is of this kind.

Lying. A lie is an utterance whose pathway is made opaque by S’s own design. It is not merely insincere, but constructed to prevent the interlocutor from recovering the causal pathway. This places it outside the free speech region by the binary prerequisite.

Lying is also a symptom of constrained expression: a speaker who lies to achieve goal G could not have achieved G through transparent expression. The lie is a workaround that reveals the constraint. This includes adversarial cases in which the goal is intrinsically incompatible with transparency, as when a fraudster claims false identity to obtain access to private information. The framework treats these symmetrically: the speaker was free to forgo G and speak transparently, or to pursue G and forgo transparency. The account does not promise that every goal a speaker has is achievable through free speech. This non-achievability is a cost to speaking freely that the lying speaker views as prohibitive.

Opinion-dependence. A consequence of grounding freedom in the cost of transparency is that whether a speaker is free to express a view depends in part on the view itself. The (T, C) cell fixes topic and context, but the cost of making the pathway transparent varies with the content of S’s mental states: two speakers in the same cell, holding different opinions on the same topic, may face very different costs of sincerity. This cuts against the ordinary picture, on which free speech is a property of the channel or the topic and indifferent to what one happens to think. We note that there is always some opinion extreme enough that voicing it carries a cost. The range of possible opinions on every topic faces a threshold beyond which sincere expression is penalised, even where more moderate views on the same topic remain free. Freedom of speech, on this account, is not uniform across the opinions a speaker might hold within a cell but falls off as the content approaches what the interlocutor(s) will not tolerate.

Intentional omission. Deliberately withholding something the speaker knows the interlocutor would need to recover the pathway is the unsaid analogue of lying: opacity is engineered through absence rather than false signal. It falls outside the free speech region by the same condition. This should be distinguished from unintentional omission: failing to mention something because it seemed obvious from context, or not thinking to include it, which involves no engineering of opacity.

5 Applications

5.1 Parrhesia and Isegoria

The causal-transparency condition asks whether a speaker’s utterance is freely produced within a given (T, C) cell, and this within-cell condition is what corresponds to Bejan’s notion of parrhesia. Since membership in $U(S, T, C)$ already guarantees the absence of Type 0-2 obstacles (section 4.1), we can define the parrhesiastic set P as the subset of $U(S, T, C)$ containing the utterances whose causal pathway is recoverable and whose cost of transparency S is prepared to bear. Its definition answers the question “what can speaker S say freely in a given cell (T, C) ?”. Whether an utterance remains in P as the cost of transparency rises is determined by a disposition of the speaker. We return to this dispositional property as willingness-to-reveal in section 5.3.

Isegoria, on the other hand, has to do with which cells can be accessed by the speaker at all. Fixing some utterance u , the set I of cells (T, C) for which u is in $U(S, T, C)$ is what we call the isegoric set. Its definition answers the question “in which cells (T, C) is it possible for S to utter u ?”. Membership in the isegoric set is fixed by the Type 0-2 conditions alone. Type 0 is always required: S must be able to form the thought that u expresses. Type 1 requires that S can produce u and determine its content, including commanding a language in which the thought can be put. Type 2 is the most salient condition for isegoria: a channel must connect S to an interlocutor who shares that language (does S have their own TV channel, or the prime minister’s email?). Type 3 plays no role here: whether S can afford to utter u frankly within a cell is a question about the parrhesiastic set of that cell, not about access to it. Keeping Type 3 out of the definition is what makes isegoria and parrhesia independent dimensions: a speaker may have access to few cells while being fully frank within them, or access to many while frank in none.

Access is not exhausted by delivery, because it is the context that confers authority on an utterance: different platforms, roles, and institutional settings give the same words different weight. A speaker may satisfy the causal-transparency condition fully within the cells of the isegoric set while lacking access to cells where the context allows their speech to carry comparable weight.

Bejan’s isegoria includes not only the right to address an assembly but to “receive a hearing” (Bejan 2020), to be engaged with, not merely delivered. The authority-to-be-listened-to component invites two readings. The first is that authority may fall under Type 2, with the “available channel” read broadly to include uptake by the receiver, so that an unread email or an unattended talk is a case of a missing channel. Alternatively, authority may be a property of the (T, C) cell itself: which cells are available to a speaker then varies with their authority. The first reading is more parsimonious, since it stays entirely within the Types 0-3 vocabulary, while the second requires an additional notion of context access. We do not require a choice between the two, and the substantive claims of this section hold under either.

5.2 Suppression Topology

Censorship, as analysed in section 4.3, is a local disruption: it raises the cost of transparency in a targeted subset of the topic-context space rather than suppressing speech as such. But suppression does not stay local. When sincere answers in adjacent cells would implicate the suppressed ones, those cells become constrained too: a speaker who cannot freely discuss the state’s economic policy without risk may be equally unable to discuss historical precedents or welfare statistics that would support the same conclusions. Suppression propagates along the inferential connections of topic-space.

The propagation is not only inferential. As McGowan’s analysis of conversational exercitives shows (section 2.3), utterances shift the norms that govern subsequent conversational moves, so suppression in one cell can also deform the conversational defaults of connected cells. Either way, the parametrisation renders the resulting structure analytically tractable: which cells a regime or platform directly suppresses, and how far the constraint spreads under inferential and normative adjacency, is a substantive question that can be posed cell by cell rather than gestured at globally. This approach involves constructing a network that maps the adjacencies between topic-context cells, then using its topology to draw conclusions about the free speech constraints on the constituent cells. The same structure is useful in analysing other applications of our definition, such as willingness-to-reveal (section 5.3).

5.3 Degree of Freedom and Willingness-to-Reveal

Section 4.2 fixed the degree to which an utterance is freely spoken as inversely proportional to the cost to S of making the causal pathway transparent. This section examines how such costs arise in practice, and distinguishes degree of freedom from a neighbouring dispositional property, willingness-to-reveal. The point that most needs illustration is that degree of freedom tracks cost, not opacity: a pathway can be perfectly recoverable while the freedom to instantiate it is near zero.

A widely reported account of Soviet electoral practice is the pure form of this situation. Take the cell (political preference, public voting), and treat a vote as an utterance in the extended sense: an expressive act that registers a preference. The utterance space contained a costless option and a costly one. Voters for the Communist party could present identification and leave; voting otherwise required entering a booth to mark the ballot. The sincere utterance was thus perfectly recoverable, since entering the booth was a maximally legible signal of dissent by design, and that is exactly where the cost was attached. The degree of freedom in the cell was near zero, not because the pathway was opaque, but because the structure of the system made transparency prohibitively expensive. No explicit prohibition, and no threat aimed at any utterance, was needed.

A related but distinct property is willingness-to-reveal: whether S would, if asked directly, clarify the causal pathway. A speaker exercising self-control who says

“I don’t want a cookie” when they do, and who would confirm this if pressed, “actually, I do, I’m just trying not to spoil my dinner”, speaks freely with high willingness-to-reveal. A speaker who would deny it if asked is exercising concealment beyond what the original utterance involves. Willingness-to-reveal is the parrhesiastic property in Bejan’s sense: the disposition to make one’s actual mental states transparent even at personal cost. As stated above, studying willingness-to-reveal is best approached through a network of contexts related to the original (T, C) cell. In particular, those in which the speaker is directly asked to clarify what was concealed. Whether S confirms or continues to conceal in such adjacent contexts is what makes the property empirically tractable. A speaker may speak freely (pathway recoverable) without being parrhesiastic. The framework keeps these distinct.

Consider a politician giving a public speech who knows about corruption in their administration. In the (general policy, public address) context, the knowledge of corruption may not surface as a communicative intention at all: it is not something they want to say, so the causal-transparency condition does not directly apply. The omission is not obviously intentional engineering of opacity, as the speaker may simply have planned to discuss other topics.

The network of related contexts makes the situation tractable. If an interviewer asks directly “were there corrupt dealings?”, this shifts the context to (corruption, direct questioning), forcing the relevant mental state to the surface. A free response would be one from which that state is recoverable. A lie, “no, nothing of the kind”, falls outside the free speech region by the prerequisite. But the characteristic response in such cases is evasion: “I am not in a position to discuss that.” Evasion is neither clearly free nor clearly lying. The interlocutor may recover from it that something is being concealed, which is a form of partial transparency about the fact of concealment, or may not, depending on context. It sits at the boundary of the recoverability condition and illustrates why studying willingness-to-reveal through adjacent contexts matters: it is in those contexts that concealment becomes visible and evasion becomes assessable.

5.4 Free Speech Law

Free speech law can be understood within the framework as fixing a specific (T, C) cell: the context in which a speaker addresses, or may be heard by, state officials, under conditions where non-legal social consequences are not at issue. A free speech law, on this reading, is a guarantee that the legal cost of transparency in that cell will not be raised above a threshold. The relation between legal protection and free speech in our sense is asymmetric. In one direction, the correspondence is total: whatever protection free speech law offers takes exactly this form, a commitment about the cost structure of specific cells. In the other direction, it is partial: not all speech that is free in our sense receives legal protection. The protected cells are a small region of the topic-context space, and even within them the law withholds protection from some utterances, lying and hate speech being the examples we treat in sections 6.3 and 6.4.

6 Objections

6.1 Freedom as a Relational Property

Freedom of speech is usually thought of as a property of the speaker, i.e. something the speaker has or lacks, independently of their audience. On our account, the same utterance can be freely spoken relative to one interlocutor and not freely spoken relative to another, depending on what contextual resources the interlocutor has available.

We take the relational character to be correct rather than objectionable. Free speech is not a property of utterances in isolation, but a property of the communicative relationship between speaker and interlocutor. This matches what we are actually interested in: a dissident's coded speech to trusted colleagues is free, but the same words in a context where the pathway to the intended audience is effectively blocked are not. A purely speaker-internal account cannot distinguish these cases, while ours can.

The relational character does not make freedom subjective or arbitrary. The shared context that determines what is recoverable is not fixed by any particular interlocutor's psychology but by the epistemic situation both parties occupy. Speakers can act to make the pathway more or less recoverable within a given context, and they can choose or shape contexts. The degree of freedom remains a fact about the communicative situation.

A related concern is how the account handles surveillance. The competent interlocutor is not any particular hearer but someone with access to the shared context of section 3. Where the speaker knows about surveillance, it enters the analysis not by changing who the interlocutor is, but by changing the context: a speaker who adapts their utterance because an informer is present is operating in a different context from one who does not, and the freedom condition (including its cost of transparency) applies to the expanded context accordingly. The case in which the speaker is unaware of surveillance falls under the misinterpretation exclusion of section 3: the speaker's model of who is reachable by the utterance is wrong, and that case is outside the framework as stated.

6.2 Degrees of Freedom

Kramer (2002) argues that freedoms do not exist by degrees: a particular freedom either exists or it does not, and gradations enter only on aggregation. If that is right, the scalar claim of section 5.3 may seem confused. The objection is answered by marking the level at which our degrees apply. The speaker's decision about a specific utterance is indeed binary: at a given cost, they either utter it or refrain, according to the threshold of cost they are prepared to bear. The degrees our account trades in do not belong to this atomic act. They belong to the cost structure across the space of utterances: as the utterance, the opinion held (section 4.3), or the context varies, the point at which the threshold is crossed moves, and it is this variation over $U(S, T, C)$ that "degree of freedom" measures.

This is also the level at which Kramer takes freedom to be measurable, since he denies gradations to particular freedoms while affirming the measurability of overall freedom (Kramer 2003; Carter 1999). Our scalar claim is a claim about the latter, and the apparent contradiction dissolves once the level is marked.

6.3 Legal Restrictions on Lying

If lying falls outside free speech on our account, an objector might conclude that laws restricting lies are not only permissible but required. This seems to license too much: government truth enforcement, mandatory sincerity, the apparatus of a thought-policing state. However, our framework characterises what free speech consists of, not what the law should protect. Even where a lie falls outside free speech on our account, the law has practical reason to reach it only when two conditions are jointly met:

- (a) The statement is provably false. Many lies cannot be established as such, as they concern contested evidence, matters of interpretation, or claims that cannot be verified. For those, the pathway may appear opaque, but we cannot establish that it is. Legal protection continues by default under epistemic uncertainty.
- (b) The statement has an identifiable victim. Where a provably false statement causes no harm to any specific party, legal intervention may not be warranted even if the statement falls outside free speech on our account. Harm is a proportionality condition on legal remedy, not a condition on whether speech is free.

These two constraints are exactly what defamation, libel, and slander law already apply. One may additionally argue that laws governing entrapment, freedom of information, screaming “fire” in a crowded theatre, and consumer protection fall under these constraints. The Brandenburg test for incitement in US constitutional law (*Brandenburg v. Ohio* 1969), which restricts speech directed at producing imminent lawless action, carves out a structurally similar region from a different angle: it targets speech whose downstream causal contribution to a foreseeable unlawful act is the source of harm, regardless of whether the pathway from intention to utterance was transparent. We return to this orthogonality in section 6.4. Our framework provides a theoretical grounding for why those carve-outs are principled: they are the practical implementation of the causal-transparency account under epistemic constraints and proportionality requirements, not arbitrary exceptions to a general liberty.

US constitutional doctrine has in fact converged on just this structure. In *United States v. Alvarez* (*United States v. Alvarez* 2012), striking down a statute that criminalised false claims of military honours, the Supreme Court held that the falsity of a statement does not by itself remove it from First Amendment protection, and that lies may be reached by law only where they produce a legally cognisable harm such as fraud, perjury, or defamation. That is exactly our two-part test read off the case law: provable falsity (a) is not sufficient for legal

reach, and an identifiable harm (b) is also required. Sunstein (Sunstein 2021) defends the same division on policy grounds, arguing that a free society must leave the great majority of falsehoods legally untouched and reserve regulation for lies that damage identifiable interests (and, as all laws, require evidence).

6.4 Hate Speech and Speech-Mediated Harm

Hate speech is a different kind of case and is orthogonal to the freedom question. Whether a speaker’s utterance is freely produced says nothing about the effects that utterance has on its targets. In speech-act-theoretic terms (Austin 1962; Searle 1969), the orthogonality is between illocutionary recoverability (our freedom condition) and perlocutionary harm (the effects produced on the target). They are independent dimensions of the same utterance.

Hate speech directed at a group can be completely free in our sense: the speaker’s intentions are transparent, the pathway recoverable, and no one is coercing them. The grounds for restricting it are not about freedom at all. Rather, they are about foreseeable harm to the target. Speech that predictably and causally raises the risk of legally quantifiable harm to identifiable others, e.g. an increased threat of violence, functions like paying a third party to cause that harm. It is an act that produces real-world consequences independently of whether the speaker’s pathway was transparent. Restricting it is not a restriction on free speech, but a limit on a category of harmful acts that happen to take a verbal form. This harm-based ground for regulating hate speech is developed at length by Waldron (Waldron 2012) and in the essays collected by Maitra and McGowan (Maitra and McGowan 2012).

Hate speech is not the only kind of speech that can cause harm independently of the freedom question. Kant’s example (Kant 1797) of the householder asked by a would-be murderer whether their target is hiding inside makes the general point: the truthful answer is freely spoken on our account (the pathway is recoverable) but causes serious foreseeable harm. The protective lie falls outside the free speech region by the binary prerequisite. The framework correctly identifies each case as freely spoken or not, but the question of which kinds of speech-mediated harm warrant restriction, and what causes some subset of these to be labelled hate speech, requires a fuller treatment of the harm side that we do not undertake here. We mark this as future work.

This completes the picture sketched in section 6.3. The speech that law restricts falls into two categories: speech that is already unfree in our sense (lying with an identifiable victim, deception producing harm), and speech restricted on grounds orthogonal to freedom (hate speech as harm to target, incitement as causal contribution to crime).

7 Conclusion and Future Work

This paper has proposed a causal-transparency account of free speech at the Type 3 level of Lowe’s taxonomy, motivated independently and extended upstream

with a Type 0 category of obstructions to thought-formation. Its central value is that it replaces a negative characterisation with a positive condition. Where the literature marks speech as unfree by the presence of threats and anticipated costs, a single structural question, “can a competent interlocutor, given an utterance u and the shared context C , recover speaker S ’s causal pathway from communicative intention to u ?”, yields as consequences what other accounts must stipulate case by case. Among the consequences we explore: (1) lying and coercion are excluded on one structural ground rather than two moral ones, (2) censorship, coercion, and chilling effects become commensurable costs on a common scale, (3) the spread of suppression through connected topics becomes a tractable structural question rather than a metaphor, (4) both the ancient vocabulary of isegoria and parrhesia and the modern practice of free speech law acquire precise content within a single definition.

The account also separates questions that free speech debates tend to run together: what freedom of speech is, which speech deserves protection, and what remedies follow. It answers the first without inheriting Mill’s empirical commitments about truth or a substantive theory of autonomy. The parametrisation by topic and context gives the answer its resolution: freedom is assessed cell by cell, so partial suppression, the normal condition of real speech environments, becomes a structure that can be mapped and studied with the tools of network analysis. Applied to law, the framework reads freedom-of-speech legislation as a guarantee about the cost of transparency in specific cells, and sorts the standard exclusions by their grounds: lying falls outside legal protection because it is not free speech to begin with, while hate speech is restricted for the harm it does, on grounds orthogonal to freedom.

The simplifying assumptions flagged in the main text mark the directions for future work: (1) partial recoverability, set aside in section 4.2; (2) the linearisation of plural and speaker-relative costs, adopted in section 3; (3) the asymmetric case in which one party misreads the shared context, as under surveillance; (4) a unified treatment of the upstream interventions at Types 0 through 2, from language persecution to infrastructure control; and (5) an account of speech-mediated harm capable of distinguishing hate speech within it. The inferential topology of topic-space stands out as a subject in its own right. That the study of freedom of speech should generate a network structure, a graph of topic-context cells along whose edges suppression propagates, is unexpected, and both its formal properties and its empirical instantiation in actual speech regimes deserve further investigation.

The causal-transparency account does not settle which speech ought to be protected. What it settles is the question that debate presupposes: what the freedom at stake is. By identifying free speech with the recoverability of the pathway from communicative intention to utterance, it hands the normative debate a defined object where it previously had intuition and precedent.

References

- Austin, J. L. 1962. *How to Do Things with Words*. Oxford University Press.
- Bejan, Teresa M. 2017. *Mere Civility: Disagreement and the Limits of Toleration*. Harvard University Press.
- Bejan, Teresa M. 2019. “Two Concepts of Freedom (of Speech).” *Proceedings of the American Philosophical Society* 163 (2): 95–107.
- Bejan, Teresa M. 2020. “Free Expression or Equal Speech?” *Social Philosophy and Policy* 37 (2): 153–69.
- Brandenburg v. Ohio*. 1969. 395 U.S. 444 (Supreme Court of the United States).
- Carter, Ian. 1999. *A Measure of Freedom*. Oxford University Press.
- Cook, Philip, and Conrad Heilmann. 2013. “Two Types of Self-Censorship: Public and Private.” *Political Studies* 61 (1): 178–96.
- Grice, H. Paul. 1957. “Meaning.” *Philosophical Review* 66 (3): 377–88.
- Grice, H. Paul. 1975. “Logic and Conversation.” In *Syntax and Semantics, Vol. 3: Speech Acts*, edited by Peter Cole and Jerry Morgan. Academic Press.
- Hornsby, Jennifer, and Rae Langton. 1998. “Free Speech and Illocution.” *Legal Theory* 4 (1): 21–37.
- Kant, Immanuel. 1797. “On a Supposed Right to Lie from Philanthropy.” In *Practical Philosophy*, edited by Mary J. Gregor. Cambridge University Press.
- Kramer, Matthew H. 2002. “Why Freedoms Do Not Exist by Degrees.” *Political Studies* 50 (2): 230–43.
- Kramer, Matthew H. 2003. *The Quality of Freedom*. Oxford University Press.
- Langton, Rae. 1993. “Speech Acts and Unspeakable Acts.” *Philosophy & Public Affairs* 22 (4): 293–330.
- Lowe, Rebecca. 2026. *Speaking Freely*. Work in progress; presented on the Elucidations Podcast, episode 155, and summarised on the author’s Substack.
- Maitra, Ishani, and Mary Kate McGowan, eds. 2012. *Speech and Harm: Controversies over Free Speech*. Oxford University Press.
- McGowan, Mary Kate. 2019. *Just Words: On Speech and Hidden Harm*. Oxford

University Press.

Mill, John Stuart. 1859. *On Liberty*. John W. Parker; Son.

Scanlon, Thomas M. 1972. "A Theory of Freedom of Expression." *Philosophy & Public Affairs* 1 (2): 204-26.

Schauer, Frederick. 1982. *Free Speech: A Philosophical Enquiry*. Cambridge University Press.

Searle, John R. 1969. *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press.

Shiffrin, Seana Valentine. 2011. "A Thinker-Based Approach to Freedom of Speech." *Constitutional Commentary* 27: 283-307.

Shiffrin, Seana Valentine. 2014. *Speech Matters: On Lying, Morality, and the Law*. Princeton University Press.

Sperber, Dan, and Deirdre Wilson. 1995. *Relevance: Communication and Cognition*. 2nd ed. Blackwell.

Sunstein, Cass R. 2021. *Liars: Falsehoods and Free Speech in an Age of Deception*. Oxford University Press.

United States v. Alvarez. 2012. 567 U.S. 709 (Supreme Court of the United States).

Waldron, Jeremy. 2012. *The Harm in Hate Speech*. Harvard University Press.